

Statistics (cdf, ppf, prf) t_test, z_test

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```


```
In [2]: t_mean = 18.7 # given
t_std = 5.8 # given
n_trips = 2000 #given
```

Question 1

```
In [3]: #Asked: should earn over 20
```

```
In [4]: import scipy.stats as stats
```

```
In [5]: x = np.linspace(t_mean-3*t_std, t_mean+3*t_std, 100)
plt.plot(x, stats.norm.pdf(x, t_mean, t_std));
plt.plot([20,20,20], [0,0.01,0.07]);
```



stats.norm.cdf

```
In [6]: #The formula to find the left area of the curve until 20
stats.norm.cdf(20, t_mean, t_std)
```

```
Out[6]: 0.5886750107430825
```

```
In [7]: # Then 1 - (above code) gives the right area of the curve
1-stats.norm.cdf(20, t_mean, t_std)
```

```
Out[7]: 0.4113249892569175
```

```
In [8]: # Now multiplying by number of trips
#(Because above calculation was for only one unit trip)
2000*0.41
```

Out [8]: 820.0

stats.binom.cdf(x, n, p)

x is the number of successes from n trials with the probability of success is p

```
In [9]: # Let's find the probability of successes of having two tails
# from three coin toss where p of having tails = 0.6
print(3*((0.6**2)*((1-0.6)**1)))
stats.binom.pmf(2,3,0.6) # pmf returns the p of successes
```

0.43199999999999994

Out [9]: 0.43199999999999994

Question 3

```
In [10]: n_trips = 5
# Asked: over 15 probability at least 3 trips out of 5 trips
right_area = 1 - stats.norm.cdf(15, t_mean, t_std)
```

```
In [11]: #"binom.cdf(num_succ-1, num_trials, prob)" Excel -> "=BINOMDIST(B2,C2,D2,1)
```

```
In [12]: print(right_area)
left_area = 1-right_area
print(left_area)
```

0.7382407127058933

0.2617592872941067

Usually dist formulas calculates left curve area so if we need to know right area, we have to 1- left_area

```
In [13]: # This formula calculates the probability of having
# 2 or less success from 5 trial whose probability = left_area
stats.binom.cdf(2,5,left_area)
```

Out [13]: 0.8836952614939727

```
In [14]: # Not having 2 or less successes from 5 trials where p = left_area
stats.binom.sf(2,5,left_area)
```

Out [14]: 0.11630473850602734

```
In [15]: #Additional
# If we are asked to find successes between [2 and 4], for example
```

```
In [16]: stats.binom.cdf(4,5,right_area)
```

Out [16]: 0.7807245754350454

In [17]:

```
stats.binom.cdf(1, 5, right_area)
```

Out[17]: 0.018557998317456286

```
In [18]: # This formula equals to excel formula --> BINOM.DIST.RANGE(5,right_area,2,4)
stats.binom.cdf(4,5,right_area)-stats.binom.cdf(1, 5, right_area)
```

Out[18]: 0.7621665771175891

Question 2

```
In [19]: n_trips = 5
# Asked: All of trip to get 15 or over
left_a = stats.norm.cdf(15, t_mean, t_std)
right_a = 1-left_a
```

```
In [20]: stats.binom.cdf(0,5,left_a)
```

Out[20]: 0.21927542456495466

Question 8

Given:

$$p = 0.6$$

$$\hat{p} = 0.5$$

$$n = 50$$

Then:

$$var = \hat{p} * (1 - \hat{p})$$

$$std = \frac{\sqrt{var}}{\sqrt{n}}$$

$$z = \frac{(p - \hat{p})}{std}$$

\hat{p} is a mean for H_a and p is from sample we got

from scipy.stats import norm

norm.cdf(1.96)

0.9750021048517795

norm.cdf(-1.96)

0.024997895148220435

stats.norm.ppf

```
In [21]: # This formula will calculate z value, two-tail
# Where confidence interval for 99 %, z value is 2.575
```

```
stats.norm.ppf(0.995)
# It means that 99% confidence interval lies
#2.575 standard deviation from the mean
#where mean = 0 and std = 1
```

Out[21]: 2.5758293035489004

```
In [22]: stats.norm.cdf(2.575)
# It will give the confidence interval of the given value
#(std of how many times far from the mean (z))
```

Out[22]: 0.9949879956682387

```
In [23]: stats.norm.cdf(1.4142)
```

Out[23]: 0.9213484060057744

```
In [24]: p_val = 1- stats.norm.cdf(1.4142)
p_val
```

Out[24]: 0.07865159399422561

Question 10

```
In [25]: m = 85
std_sample = 32
n = 16
```

```
In [26]: std_r = std_sample/np.sqrt(n)
```

```
In [27]: z =stats.norm.ppf(0.95)
```

```
In [28]: marg_errors = z*std_r
```

```
In [29]: conf_int_min = m-marg_errors
conf_int_max = m+marg_errors
```

```
In [30]: print(conf_int_min, 'to ', conf_int_max)
```

71.84117098438823 to 98.15882901561177

```
In [31]: #Studnt, n=999, p<0.05, 2-tail\
#equivalent to Excel T.INV(0.05,999) or T.INV.2T(0.1, 999)
#90% confidence interval
#print stats.t.ppf(1-0.025, 999)

#Studnt, n=999, p<0.05%, Single tail\
#equivalent to Excel TINV(2*0.05,999)\
#print stats.t.ppf(1-0.05, 999)
```

In [32]: `#But t is used more`

stats.t.ppf

```
In [33]: t = stats.t.ppf(0.95,15)
marg_errors = std_r*t
conf_int_min = m-marg_errors
conf_int_max = m+marg_errors
print(conf_int_min, 'to ', conf_int_max)
```

70.97559715445956 to 99.02440284554044

Question 14

```
In [34]: uber = pd.read_excel('data2.xlsx', sheet_name = 'Uber pay')
```

```
In [35]: uber.head(2)
```

```
Out[35]:
```

	DriverID	City	Gender	Trips	Hourly Pay	Unnamed: 5
0	1	New York	Male	493	7.49	NaN
1	2	Boston	Male	1195	8.20	NaN

```
In [36]: chic_mal = uber[(uber['City']=='Chicago') & (uber['Gender']=='Male')]['Hourly Pay']
```

```
In [37]: std = chic_mal.std()
mean = chic_mal.mean()
n = len(chic_mal)
```

```
In [38]: t = stats.t.ppf(0.975, 181)
std_n = std/np.sqrt(n)
```

```
In [39]: marg_err = t*std_n
marg_err
```

```
Out[39]: 0.27923413996605523
```

Question 15

```
In [40]: fem_pay = uber[uber['Gender']=='Female']['Hourly Pay']
```

```
In [41]: mean = fem_pay.mean()
std = fem_pay.std()
n = len(fem_pay)
```

```
In [42]: t = stats.t.ppf(0.975, n-1)
std_n = std/np.sqrt(n)
marg_err = t*std_n
conf_int_min = mean-marg_err
```

```
conf_int_max = mean+marg_err  
print(conf_int_min, conf_int_max)
```

7.211959435480943 7.670625598532671

```
In [43]: male = uber['Gender'].value_counts()[0]  
female = uber['Gender'].value_counts()[1]
```

```
In [44]: p = female/(male+female)  
n = male+female
```

```
In [45]: z = stats.norm.ppf(0.975)
```

```
In [46]: t = stats.t.ppf(0.975, n-1)
```

```
In [47]: std = np.sqrt(p*(1-p)/n)
```

```
In [48]: marg_err = t*std  
max_conf = p+marg_err  
min_conf = p-marg_err  
print(min_conf, max_conf)
```

0.2657283864579868 0.32227161354201317

Question 16

```
In [49]: # Ho: m>= m_sample  
# Ha: m< m_sample
```

```
In [50]: population_mean = 7.25  
sample_mean = uber['Hourly Pay'].mean()  
alfa = 0.05
```

```
In [51]: n = len(uber)  
samp_std = uber['Hourly Pay'].std()  
pop_std = samp_std/np.sqrt(n)
```

```
In [52]: z = (sample_mean - population_mean)/pop_std
```

```
In [53]: print(sample_mean)  
print(samp_std)  
n
```

7.8934099999999993
1.9947413528989075
1000

Out[53]:

stats.t.cdf

In [54]:

```
p = 1 - stats.t.cdf(z, 999, 0)
p
# we will reject the null H and agree with the Ha
# that uber driver earns more than pop_mean
```

Out[54]: 0.0

In [55]: pop_std

Out[55]: 0.06307926018086264

In [56]: stats.norm.cdf(7.89, 7.25, 0.06)

Out[56]: 1.0

Question 18

Two sample t test

In [57]:
chicago = uber[uber['City'] == 'Chicago']['Hourly Pay']
york = uber[uber['City'] == 'New York']['Hourly Pay']

In [58]:
ch_m = chicago.mean()
ch_std = chicago.std()
ch_n = len(chicago)
york_m = york.mean()
york_std = york.std()
york_n = len(york)

$$t = \frac{m1 - m2}{\sqrt{\frac{s1^2}{n1} + \frac{s2^2}{n2}}}$$

In [59]: print(ch_m, ch_std, ch_n, york_m, york_std, york_n)

8.071140684410652 1.998541508080747 263 7.882751842751838 1.9597766271731287 407

In [60]:
Ho: ch_m <= york_m
Ha: ch_m > york_m
t = (ch_m - york_m) / np.sqrt((ch_std**2)/ch_n + (york_std**2)/york_n)

In [61]: stats.t.cdf(t, ch_n+york_n-2)

Out[61]: 0.8848237669591617

In [62]:
just another example:

In [63]:
#Blackboard: Given:
m1 = 82
s1 = 2.4

```
n1 = 15  
  
m2 = 84  
s2 = 1.7  
n2 = 12  
alfa = 0.05
```

```
In [64]: t = (m1-m2)/np.sqrt((s1**2)/n1 + (s2**2)/n2)  
t
```

```
Out[64]: -2.5301595052287063
```

```
In [65]: p = stats.t.cdf(t, 25)
```

```
In [66]: # It is two tail test, because m1 != m2 (either high or low)  
# So alfa = alfa/2 = 0.025  
print(p)  
alfa/2
```

```
Out[66]: 0.00903684801412631  
0.025
```

****p is lower than alfa, so we reject Ho and accept Ha, which is:**

there is significant difference between mean1 and mean2**

```
In [67]: hotel = pd.read_excel('data2.xlsx', sheet_name = 'hotel')
```

```
In [68]: hotel.head()
```

```
Out[68]:
```

	BookingID	Channel	Weekday	Rate	Cancelled
0	1	Online	Tuesday	385.56	False
1	2	Online	Wednesday	305.15	False
2	3	Online	Thursday	559.00	True
3	4	Online	Monday	424.15	False
4	5	Online	Sunday	459.00	True

Question 19

```
In [69]: rate = hotel['Rate']
```

```
In [70]: r_m = rate.mean()  
r_s = rate.std()  
r_n = len(rate)
```

```
In [71]: # Asked: build 95% confidence interval for rate  
z = stats.norm.ppf(0.975)  
t = stats.t.ppf(0.975, 999)  
print(z, t)
```


1.959963984540054 1.9623414611334487

```
In [72]: pop_std = r_s/np.sqrt(r_n)
```

```
In [73]: marg_er = pop_std*z
```

```
In [74]: min_conf = r_m -marg_er
max_conf = r_m + marg_er
print(min_conf, max_conf)

295.33095035193713 304.31864679106235
```

```
In [75]: marg_er = pop_std*t
min_conf = r_m -marg_er
max_conf = r_m + marg_er
print(min_conf, max_conf)

295.3254992216916 304.3240979213079
```

Question 20

```
In [76]: # alfa = 0.05
# Ho: r_m <= 300
# Ha: r_m > 300
```

```
In [77]: t = (r_m-300)/pop_std # We should (x-x_mean)/(std/sqrt(n))
```

```
In [78]: p = stats.t.cdf(t, 999) # to find how much area it will
# cover when the distance from mean is equal to t*std
```

```
In [79]: p
```

```
Out[79]: 0.4695529077738393
```

Question 21

```
In [80]: online = hotel[hotel['Channel']=='Online']['Rate']
offline = hotel[hotel['Channel']!='Online']['Rate']
```

```
In [81]: of_mean = offline.mean()
of_std = offline.std()
of_len = len(offline)
on_mean = online.mean()
on_std = online.std()
on_len = len(online)
```

```
In [82]: # Ho: of_mean >= on_mean
# Ha: of_mean < on_mean
alpha = 0.05
```

```
In [83]:
```

```
t = (of_mean-on_mean)/np.sqrt(of_std**2/of_len + on_std**2/on_len)
```

```
In [84]: stats.t.cdf(t, of_len+on_len-2) # p value is higher than alpha
```

```
Out[84]: 0.2890174472520236
```

Question 22

```
In [85]: hotel[hotel['Channel']=='Online'][['Rate',  
                                           'Cancelled']].groupby('Cancelled').count()
```

```
Out[85]:
```

	Rate
Cancelled	
False	611
True	224

```
In [86]: cancelled = 224  
not_cancelled = 611  
percent = 20  
# Let's start  
p_pop = 0.20  
p_sample = cancelled/(cancelled+not_cancelled)  
n = cancelled+not_cancelled  
std = np.sqrt(p_pop*(1-p_pop)/n)  
z = (p_sample - p_pop)/std  
z
```

```
Out[86]: 4.9314158489676245
```

```
In [87]: # Ho: p_sample <= 0.2  
# Ha: p_sample > 0.2  
alpha = 0.05  
upper_bound = 0.95
```

```
In [88]: print(p_sample, std)
```

```
0.2682634730538922 0.013842570804119656
```

```
In [89]: p = stats.norm.cdf(z)
```

```
In [90]: # stats.binom.cdf(p_sample, 0.2, std) is different thing,  
# we can't use. (above)  
stats.norm.cdf(p_sample*100, p_pop*100, std*100)
```

```
Out[90]: 0.9999995918213253
```

```
In [91]: p # it is higher than upper bound, so we reject null H,  
# therefore p_sample is higher than 0.2
```

```
Out[91]: 0.9999995918213253
```

Question 23

```
In [92]: stats.norm.cdf(6,8,12) # area until 6 so 1- ..  
# gives the area starting from 6 till end
```

```
Out[92]: 0.43381616738909634
```

```
In [93]: stats.norm.cdf(6,15,30) # the same like above
```

```
Out[93]: 0.3820885778110474
```

Question 24

```
In [94]: 1-stats.norm.cdf(40,32.4, 3.6) # over 40
```

```
Out[94]: 0.017381381311111422
```

Question 25

```
In [95]: stats.norm.ppf(0.75, 32.4, 3.6)  
# gives the percentile (counter cdf)  
# it gives 75% point where mean = 32.4 and std = 3.6
```

```
Out[95]: 34.82816310070589
```

Question 26

```
In [96]: z_m = 8  
z_s = 12  
c_m = 15  
c_s = 30
```

```
In [97]: m_diff = z_m-c_m  
s_diff = np.sqrt(z_s**2 + c_s**2)  
z = m_diff/s_diff  
#Asked: when m_diff < 0
```

```
In [98]: m_diff
```

```
Out[98]: -7
```

```
In [99]: stats.norm.cdf(0, m_diff, s_diff) # it is when m_diff < 0
```

```
Out[99]: 0.5857573222326861
```

```
In [100]: stats.norm.cdf(z)
```

```
Out[100]: 0.41424267776731394
```

Question 27

```
In [101...  
n = 8  
price = 30  
mean = 0.08  
std = 0.12  
total = 400
```

```
In [102...  
single = total/n  
single
```

```
Out[102... 50.0
```

```
In [103...  
stats.norm.cdf(50, 32.4, 3.6) # till 50  
1-stats.norm.cdf(50, 32.4, 3.6) # over 50  
# almost 0
```

```
Out[103... 5.070335422630023e-07
```

```
In [104...  
c_price = 30  
c_mean= 0.08  
c_std = 0.12  
z_price = 40  
z_mean = 0.15  
z_std = 0.3  
minimize = 50  
n_share = 20
```

```
In [105...  
stats.norm.ppf(0.5, 32.4, 3.6)
```

```
Out[105... 32.4
```

```
In [106...  
new_df=pd.DataFrame(np.arange(1,21))
```

```
In [107...  
new_df['zoom_price'] = z_price  
new_df['zoom_mean'] = z_mean  
new_df['zoom_std'] = z_std  
new_df['cloud_price'] = c_price  
new_df['cloud_mean'] = c_mean  
new_df['cloud_std'] = c_std
```

```
In [108...  
new_df = new_df.rename(columns= {0:'share_count'})
```

```
In [109...  
new_df['zoom_return'] = new_df['share_count']*(new_df['zoom_price']+  
                                              new_df['zoom_price']*new_df['zoom_mean'])
```

```
In [110...  
new_df['zoom_return_std'] = np.sqrt(new_df['zoom_std']**2*new_df['share_count'])
```

```
In [111...  
new_df['cloud_return'] = new_df['share_count']*(new_df['cloud_price']+  
                                              new_df['cloud_price']*new_df['cloud_mean'])  
new_df['cloud_return_std'] = np.sqrt(new_df['cloud_std']**2*new_df['share_count'])
```

In [112...

new_df

Out[112...

	share_count	zoom_price	zoom_mean	zoom_std	cloud_price	cloud_mean	cloud_std	zoom_return	zo
0	1	40	0.15	0.3	30	0.08	0.12	46.0	
1	2	40	0.15	0.3	30	0.08	0.12	92.0	
2	3	40	0.15	0.3	30	0.08	0.12	138.0	
3	4	40	0.15	0.3	30	0.08	0.12	184.0	
4	5	40	0.15	0.3	30	0.08	0.12	230.0	
5	6	40	0.15	0.3	30	0.08	0.12	276.0	
6	7	40	0.15	0.3	30	0.08	0.12	322.0	
7	8	40	0.15	0.3	30	0.08	0.12	368.0	
8	9	40	0.15	0.3	30	0.08	0.12	414.0	
9	10	40	0.15	0.3	30	0.08	0.12	460.0	
10	11	40	0.15	0.3	30	0.08	0.12	506.0	
11	12	40	0.15	0.3	30	0.08	0.12	552.0	
12	13	40	0.15	0.3	30	0.08	0.12	598.0	
13	14	40	0.15	0.3	30	0.08	0.12	644.0	
14	15	40	0.15	0.3	30	0.08	0.12	690.0	
15	16	40	0.15	0.3	30	0.08	0.12	736.0	
16	17	40	0.15	0.3	30	0.08	0.12	782.0	
17	18	40	0.15	0.3	30	0.08	0.12	828.0	
18	19	40	0.15	0.3	30	0.08	0.12	874.0	
19	20	40	0.15	0.3	30	0.08	0.12	920.0	

In [113...

drive = pd.read_excel('data2.xlsx',
sheet_name= 'drive-through')

In [114...

drive

Out[114...

	Customer	Time
0	1	138.6
1	2	43.3
2	3	57.5
3	4	84.2
4	5	44.7
...
195	196	88.4
196	197	198.4
197	198	36.9
198	199	47.6

	Customer	Time
199	200	50.7

200 rows × 2 columns

Question 29

```
In [115... st_err = stats.sem(drive['Time']) # standard error
st_err
```

```
Out[115... 4.7617343188819445
```

```
In [116... drive['Time'].std()/np.sqrt(len(drive))
# equivalent for std_error = std/sqrt(n)
```

```
Out[116... 4.7617343188819445
```

1 st way for confidence interval when alpha = 0.90

```
In [117... mean= drive['Time'].mean()
```

```
In [118... t = stats.t.ppf(0.95, len(drive) - 1)
```

```
In [119... marg_err = t*st_err
```

```
In [120... mean-marg_err
```

```
Out[120... 89.31801144522497
```

```
In [121... mean+marg_err
```

```
Out[121... 105.05598855477506
```

2 nd way for confidence interval when alpha = 0.90

```
In [122... stats.t.interval(alpha = 0.9,
                    df = len(drive)- 1,
                    loc = np.mean(drive['Time']),
                    scale = stats.sem(drive['Time']))
```

```
Out[122... (89.31801144522497, 105.05598855477506)
```

Question 30

```
In [123... light = pd.read_excel('data2.xlsx', sheet_name = 'lightbulbs')
```

```
In [124... light.head()
```

Out [124...	Lightbulb	Lifetime
	0	1 840.08
	1	2 960.00
	2	3 953.38
	3	4 981.14
	4	5 938.66

```
In [125...  
# Ho: life <= 1000;  
# Ha: life > 1000;  
alpha = 0.05
```

1st way

```
In [126...  
mean = light['Lifetime'].mean()  
std = light['Lifetime'].std()  
n = len(light)  
pop_std = std/np.sqrt(n)
```

```
In [127...  
t = (mean -1000)/pop_std  
t
```

```
Out[127...  
-1.0303562006210676
```

```
In [128...  
stats.t.cdf(t, n-1)  
# this tells us that  
# sample_mean - pop_mean < 0 has the area of 0.15,  
# so 1- p gives the area of sample_mean-pop_mean > 0  
1- stats.t.cdf(t, n-1) # this is final p value
```

```
Out[128...  
0.8473232451095432
```

second way

```
In [129...  
stats.ttest_1samp(light['Lifetime'],  
                  1000,  
                  alternative = 'greater')
```

```
Out[129...  
Ttest_1sampResult(statistic=-1.0303562006210678, pvalue=0.8473232451095432)
```

Question 31:

```
In [130...  
battery = pd.read_excel('data2.xlsx', sheet_name = 'battery')
```

```
In [131...  
# Ho: brand_1 = brand2  
# Ha: brand1 != brand2  
alpha = 0.05
```

1st:

```
In [132... battery.columns = ['battery', 'brand_1', 'brand_2']
```

```
In [133... b1_mean=battery['brand_1'].mean()
b1_std = battery['brand_1'].std()
b2_mean=battery['brand_2'].mean()
b2_std = battery['brand_2'].std()
b1_len = len(battery['brand_1'])
b2_len = len(battery['brand_2'])
```

```
In [134... t = (b1_mean - b2_mean)/np.sqrt(b1_std**2/ b1_len + b2_std**2/b2_len)
```

```
In [135... t
```

```
Out[135... 0.7677462249616107
```

```
In [136... print(1 -stats.t.cdf(t, b1_len +b2_len -2))
# alpha = 0.05/0.95 so we won't reject the null H;
# There is no signifiant difference
print((1 -stats.t.cdf(t, b1_len +b2_len -2))*2)
```

```
0.22177628637630398
0.44355257275260795
```

2nd:

```
In [137... stats.ttest_ind(battery['brand_1'],
                  battery['brand_2'],
                  alternative = 'two-sided')
# p value is 0.46, so divided by two (two_sided)
```

```
Out[137... Ttest_indResult(statistic=0.76774622496159, pvalue=0.4435525727526203)
```

```
In [138... df = pd.DataFrame({'col_1':[10, 9, 11, 8, 14, 6],
                    'col_2':[11, 12, 7, 10, 8,13]})
```

```
In [139... mean_1 = df['col_1'].mean()
std_1 = df['col_1'].std()
n_1 = len(df)
mean_2 = df['col_2'].mean()
std_2 = df['col_2'].std()
n_2 = len(df)
```

```
In [140... t = (mean_1 - mean_2)/np.sqrt(std_1**2/n_1 + std_2**2/n_2)
p = stats.t.cdf(t, n_1+n_2-2)
print(t, p)
```

```
-0.34188172937891387 0.3697579336002253
```

```
In [141... stats.ttest_ind(df['col_1'],
                  df['col_2'],
                  alternative='two-sided')
```

```
Out[141... Ttest_indResult(statistic=-0.34188172937891387, pvalue=0.7395158672004506)
```


Question 33

```
In [142... train = pd.read_excel('data2.xlsx', sheet_name = 'training')
train.head(3)
```

```
Out[142...      Employee  Score_Before  Score_After
0           1           84           84
1           2           66           85
2           3           84           87
```

```
In [143... before = train['Score_Before']
after = train['Score_After']
```

1st:

```
In [144... b_m = before.mean()
b_s = before.std()
a_m = after.mean()
a_s = after.std()
n = len(before)
```

```
In [145... t = (b_m - a_m)/np.sqrt(b_s**2/n + a_s**2/n)
```

```
In [146... t
```

```
Out[146... -7.241961961405006
```

```
In [147... 1-stats.t.cdf(t, 2*n-2)
```

```
Out[147... 0.9999999998640184
```

2nd

```
In [148... stats.ttest_ind(before, after, alternative = 'greater')
```

```
Out[148... Ttest_indResult(statistic=-7.241961961405005, pvalue=0.9999999998640184)
```

```
In [149... !jupyter nbconvert --to webpdf Assign_cdf.ipynb
```

```
[NbConvertApp] Converting notebook Assign_cdf.ipynb to webpdf
[NbConvertApp] Writing 474545 bytes to Assign_cdf.pdf
```